

## **Database Tip #2 -- How the Optimizer Uses Indexes**


When a query includes a filter, the optimizer must choose the most efficient way to retrieve the required rows. That choice almost always comes down to two access paths:

- Use an index seek to jump directly to matching rows, or
- Perform a table or index scan to read pages sequentially.

The optimizer makes this decision using a simple but powerful concept: selectivity — how many rows the filter is expected to return.

To understand how this works, here's a clear, step-by-step decision flow that mirrors the optimizer's logic.

## Optimizer Decision Flow: Seek or Scan?



**Note:** This diagram shows the optimizer's behavior when evaluating a single filter condition. With multiple predicates, the optimizer considers each indexed column, their combined selectivity, and whether composite indexes exist.

## **Understanding Selectivity**

Selectivity determines whether the index is helpful:

- High selectivity: the filter returns a small percentage of the table
- Low selectivity: the filter returns a large percentage of the table

And importantly...if the query has no filter at all, the optimizer must read the entire table — which always results in a scan, regardless of the table's size.

On small tables, this is fast.

On large tables, it can be expensive — but still cheaper than using an index when all rows are needed.

## **Why Scans Aren't Always "Bad"**

A scan simply means: "Read all data pages sequentially in storage order."

Sequential I/O is extremely efficient as the DBMS reads large amounts of data with minimal overhead, which is why scans can outperform seeks when many rows are needed.

And here's the key nuance: whether a scan is "fast" or "slow" is relative to the size of the table.

- On small tables, scanning every page is often faster than seeking into an index.
- On large tables, a scan can be expensive — but still cheaper than performing thousands or millions of random lookups when many rows are needed.

Indexes help only when they allow the optimizer to avoid reading non-qualifying data. If most of the table must be read anyway, the index provides no benefit.

## **Takeaway**

Indexes improve performance by reducing the amount of data the DBMS must read.

If the filter is highly selective, a seek is efficient.

If the filter returns many rows — or there's no filter at all — a scan becomes the cheapest option.